

# Predicting quantity of sales given a price

*Florian Peixoto*

The main objective of our model will be predict the quantity that will be sold, given a determined price to a product.

Frist of all let's load the data we have.

```
sales <- read.csv2("sales.csv", sep=";", stringsAsFactors = T)
prices <- read.csv2("comp_prices.csv", sep=";", stringsAsFactors = F)
```

As we can see the data we want is divided between the both sets, the quantity is in sales and the prices is in comp\_prices, so we need to find a way to merge the two sets.

```
names(sales)
```

```
## [1] "PROD_ID"      "DATE_ORDER"  "QTY_ORDER"   "REVENUE"
```

```
names(prices)
```

```
## [1] "PROD_ID"      "DATE_EXTRACTION"  "COMPETITOR"
## [4] "COMPETITOR_PRICE" "PAY_TYPE"
```

Between the two sets we can find PROD\_ID and a date field, we can use them to merge it to a unique set, but the problem is that we have a lot of prices listed to the same date, two for each competitor. So we can transform this field into derivated fields, let's try to use min, max, mean, median and standard deviation. The data field could be transformed as well, so let's create a field for each attribute.

```
prices_processed <- filter(prices, complete.cases(prices)) %>%
  mutate(YEAR = lubridate::year(as_datetime(DATE_EXTRACTION)),
         MONTH = lubridate::month(as_datetime(DATE_EXTRACTION)),
         DAY = lubridate::day(as_datetime(DATE_EXTRACTION)),
         COMPETITOR_PRICE = as.numeric(COMPETITOR_PRICE),
         PROD_ID = as.factor(PROD_ID)) %>%
  group_by(PROD_ID, YEAR, MONTH, DAY) %>%
  summarise(MIN_PRICE = min(COMPETITOR_PRICE),
            MAX_PRICE = max(COMPETITOR_PRICE),
            MEAN_PRICE = mean(COMPETITOR_PRICE),
            SD_PRICE = sd(COMPETITOR_PRICE),
            MEDIAN_PRICE = median(COMPETITOR_PRICE))
```

```
prices_processed
```

```
## Source: local data frame [1,929 x 9]
```

```
## Groups: PROD_ID, YEAR, MONTH [?]
```

```
##
```

```
##   PROD_ID  YEAR MONTH  DAY MIN_PRICE MAX_PRICE MEAN_PRICE SD_PRICE
##   <fctr> <dbl> <dbl> <int>    <dbl>    <dbl>    <dbl>    <dbl>
## 1      P1  2015     3    15   1499.00     1499   1499.000   0.00000
## 2      P1  2015     3    16   1362.50     1499   1464.205  58.47609
## 3      P1  2015     3    17   1362.50     1499   1429.793  65.61760
## 4      P1  2015     3    18   1362.50     1499   1441.008  64.24923
## 5      P1  2015     3    19   1304.13     1499   1407.363  80.50857
## 6      P1  2015     3    20   1304.13     1499   1439.727  74.83624
## 7      P1  2015     3    21   1359.00     1499   1439.842  65.80861
```

```
## 8      P1 2015      3    22 1424.05      1499 1488.293 27.21708
## 9      P1 2015      3    23 1403.90      1499 1456.487 46.08149
## 10     P1 2015      3    24 1403.90      1499 1454.808 46.68677
## # ... with 1,919 more rows, and 1 more variables: MEDIAN_PRICE <dbl>
```

Now let's process the sales data too.

```
sales_processed <- filter(sales, complete.cases(sales)) %>%
  mutate(YEAR = lubridate::year(as_datetime(DATE_ORDER)),
         MONTH = lubridate::month(as_datetime(DATE_ORDER)),
         DAY = lubridate::day(as_datetime(DATE_ORDER)),
         QTY_ORDER = as.numeric(QTY_ORDER),
         PROD_ID = as.factor(PROD_ID)) %>%
  group_by(PROD_ID, YEAR, MONTH, DAY) %>%
  summarise(QTY_ORDER = sum(QTY_ORDER))
```

```
sales_processed
```

```
## Source: local data frame [2,162 x 5]
## Groups: PROD_ID, YEAR, MONTH [?]
##
##   PROD_ID YEAR MONTH DAY QTY_ORDER
##   <fctr> <dbl> <dbl> <int>    <dbl>
## 1      P1 2015      2     4        10
## 2      P1 2015      2     5        12
## 3      P1 2015      2     6        21
## 4      P1 2015      2     7         4
## 5      P1 2015      2     8         7
## 6      P1 2015      2     9         5
## 7      P1 2015      2    10        10
## 8      P1 2015      2    11        11
## 9      P1 2015      2    12        16
## 10     P1 2015      2    13         7
## # ... with 2,152 more rows
```

Now we can merge the two data sets.

```
merged_data <- merge(sales_processed, prices_processed)

tail(merged_data)
```

```
##   PROD_ID YEAR MONTH DAY QTY_ORDER MIN_PRICE MAX_PRICE MEAN_PRICE
## 1913     P9 2015      9   4         44   429.90      579   454.3022
## 1914     P9 2015      9   5         42   411.48      579   453.7212
## 1915     P9 2015      9   6         39   411.48      579   452.2540
## 1916     P9 2015      9   7        118   411.48      579   452.6156
## 1917     P9 2015      9   8         75   411.48      579   452.3818
## 1918     P9 2015      9   9         55   411.48      579   451.4950
##   SD_PRICE MEDIAN_PRICE
## 1913 45.10212      429.90
## 1914 47.13315      429.90
## 1915 44.00008      431.52
## 1916 46.43189      429.90
## 1917 48.13763      429.90
## 1918 44.58876      433.14
```

Now we'll divide the data into two sets, the training set and the test set, we use this to ensure a trustworthy model

validated by results we previously know. We'll use 80% of the data to the training set and the rest to the test set.

```
merged_data.hex <- as.h2o(merged_data, destination_frame = "merged_data.hex")
```

```
##
|
|
|
|=====| 100%
```

```
data_to_model <- h2o.splitFrame(data = merged_data.hex ,
                                ratios = 0.80,
                                seed=1234)
```

```
train <- data_to_model[[1]]
test <- data_to_model[[2]]
```

Now we'll try to find a good linear regression model that can predicts our price. As our target, we'll use the MEDIAN\_PRICE as it's a real value of price present in our data.

For this work we'll use (definitions from: [https://github.com/h2oai/h2o-training-book/blob/master/hands-on\\_training/regression.md](https://github.com/h2oai/h2o-training-book/blob/master/hands-on_training/regression.md)):

- Generalized Linear Models (GLM): Average an ensemble of weakly predicting (small) trees where each tree “adjusts” to the “mistakes” of the preceding trees.
- Gradient (Tree) Boosting Machines (GBM): Average an ensemble of weakly predicting (small) trees where each tree “adjusts” to the “mistakes” of the preceding trees.
- Random Forests: Average an ensemble of weakly predicting (larger) trees where each tree is de-correlated from all other trees.

```
result <- c()
```

```
model <- c("GBM", "GLM" , "RF")
```

```
fields <- setdiff(setdiff(names(train), "QTY_ORDER"), "MEDIAN_PRICE")
```

```
GBM <- h2o.gbm(x = fields, build_tree_one_node = T,
               y = "QTY_ORDER",
               training_frame = train,
               validation_frame = test,
               seed=1234)
```

```
##
|
|
|
|=====| 44%
|
|=====| 100%
```

```
result[[1]] <- h2o.r2(GBM, valid = TRUE)
```

```
GLM <- h2o.glm(x = fields,
                y = "QTY_ORDER",
                training_frame = train,
```

```

validation_frame = test,
family = "poisson",
seed=1234)

##
|
|
|
|=====| 100%

result[[2]] <- h2o.r2(GLM, valid = TRUE)

RF <- h2o.randomForest(x = fields,
y = "QTY_ORDER",
training_frame = train,
validation_frame = test,
seed=1234)

##
|
|
|
|=====| 16%
|
|=====| 100%

result[[3]] <- h2o.r2(RF, valid = T)

results <- data.frame(model,result)


names(results) <- c("model", "squared_ratios")

results

## model squared_ratios
## 1 GBM 0.2927767
## 2 GLM 0.2779196
## 3 RF 0.3995000

ggplot(data=results, aes(x = model, y = squared_ratios, fill = model)) +
geom_bar(stat="identity", position = "dodge")

```



As we can see the random forest had the best score above all others, although the rate was not too impressive, so let's use it and see what more this model has to tell us.

```
RF_summary <- summary(RF)
```

```
## Model Details:
## =====
##
## H2ORegressionModel: drf
## Model Key: DRF_model_R_1477243664145_252
## Model Summary:
##   number_of_trees number_of_internal_trees model_size_in_bytes min_depth
## 1           50           50           431193           19
##   max_depth mean_depth min_leaves max_leaves mean_leaves
## 1         20  19.94000       294       849    681.30000
##
## H2ORegressionMetrics: drf
## ** Reported on training data. **
## ** Metrics reported on Out-Of-Bag training samples **
##
## MSE:  339102.6
## RMSE:  582.3252
## MAE:  226.5663
## RMSLE:  1.166771
## Mean Residual Deviance :  339102.6
##
##
```


```

## H2ORegressionMetrics: drf
## ** Reported on validation data. **
##
## MSE: 228748
## RMSE: 478.2761
## MAE: 195.3498
## RMSLE: 1.130537
## Mean Residual Deviance : 228748
##
##
##
## Scoring History:
##      timestamp      duration number_of_trees training_rmse
## 1 2016-10-23 20:59:57 0.003 sec                0
## 2 2016-10-23 20:59:57 0.005 sec                1      804.38679
## 3 2016-10-23 20:59:57 0.021 sec                2      751.94285
## 4 2016-10-23 20:59:57 0.029 sec                3      678.41282
## 5 2016-10-23 20:59:57 0.036 sec                4      662.39115
##      training_mae training_deviance validation_rmse validation_mae
## 1
## 2      291.80519      647038.10250      624.14024      223.87314
## 3      273.17998      565418.05247      510.94162      186.17343
## 4      251.93005      460243.94775      493.22184      185.16384
## 5      245.43263      438762.03954      486.20123      183.31704
##      validation_deviance
## 1
## 2      389551.03933
## 3      261061.34282
## 4      243267.78820
## 5      236391.63733
##
## ---
##      timestamp      duration number_of_trees training_rmse
## 46 2016-10-23 20:59:58 0.506 sec                45      581.01261
## 47 2016-10-23 20:59:58 0.522 sec                46      581.08096
## 48 2016-10-23 20:59:58 0.538 sec                47      583.00305
## 49 2016-10-23 20:59:58 0.549 sec                48      583.14751
## 50 2016-10-23 20:59:58 0.564 sec                49      584.46314
## 51 2016-10-23 20:59:58 0.580 sec                50      582.32520
##      training_mae training_deviance validation_rmse validation_mae
## 46      227.48140      337575.65440      479.94194      195.79442
## 47      227.30703      337655.07915      479.19784      195.26834
## 48      227.26138      339892.56157      478.46123      195.47823
## 49      226.88415      340061.01890      479.27009      195.00290
## 50      227.37359      341597.16733      477.78017      194.88105
## 51      226.56632      339102.64393      478.27606      195.34978
##      validation_deviance
## 46      230344.26112
## 47      229630.56557
## 48      228925.14389
## 49      229699.81772
## 50      228273.89458
## 51      228747.98951

```

```
##
## Variable Importances: (Extract with `h2o.varimp`)
## =====
##
## Variable Importances:
##   variable relative_importance scaled_importance percentage
## 1  PROD_ID 5758912512.000000      1.000000    0.226696
## 2  MIN_PRICE 4524772352.000000      0.785699    0.178115
## 3  MEAN_PRICE 4128257792.000000      0.716847    0.162506
## 4    DAY 3647966464.000000      0.633447    0.143600
## 5  SD_PRICE 2659613696.000000      0.461826    0.104694
## 6  MAX_PRICE 2469554176.000000      0.428823    0.097212
## 7    MONTH 2214638592.000000      0.384558    0.087178

g <- ggplot(data = RF_summary, aes(x = variable, y = scaled_importance, fill = variable)) +
  geom_bar(stat = "identity", position = "dodge")
g
```



We found out that the field YEAR was irrelevant for the set, so we can remove it, but for meanings of further inspection we'll let it be. We could detect the principal fields within the summary of the model:

- PROD\_ID
- DAY
- MEAN\_PRICE
- MIN\_PRICE

It appears that some products sell more than others and some days are more important as well, probably it has something to do with the day of week. The mean and min could work together as lower prices would decrease the mean and make the product more affordable.

Let's see if we can make it even more precise.

```
RF2 <- h2o.randomForest(
  x = fields,
  y = "QTY_ORDER",
  training_frame = train,
  validation_frame = test,
  ntrees = 75,
  max_depth = 35,
  seed = 1234
)
```

```
## Warning in .h2o.startModelJob(algo, params, h2oRestApiVersion): Dropping constant columns: [YEAR].
```

```
##
|
|
|
|=====| 12%
|
|=====| 93%
|
|=====| 100%
```

```
h2o.r2(RF2, valid = T)
```

```
## [1] 0.4140149
```

It represents a small increase in precision but it's not a good one, let's try to create a more accurate model. We could use the DAY variable, let's find the median prices for each day and what day of week it is. The week must have something to do with the sales as well, so we'll calculate the median

```
merged_data <-
  mutate(merged_data,
    WD = wday(ymd(paste(
      YEAR, sprintf("%02d", MONTH), sprintf("%02d", DAY), sep = ""))),
    WEEK = lubridate::week(ymd(paste(
      YEAR, sprintf("%02d", MONTH), sprintf("%02d", DAY), sep = ""))))

mean_prices_wd <- group_by(merged_data, WD) %>%
  summarise(MEAN_WD = median(MIN_PRICE))

mean_prices_week <- group_by(merged_data, WEEK) %>%
  summarise(MEAN_WEEK = median(MIN_PRICE))

merged_data <- merge(merged_data, mean_prices_wd)

merged_data <- mutate(merged_data,
  DIF_WD = MEDIAN_PRICE / MEAN_WD,
  DIF_WEEK = MEDIAN_PRICE / MEAN_WEEK)

tail(merged_data)
```



```
##      WD PROD_ID YEAR MONTH DAY QTY_ORDER MIN_PRICE MAX_PRICE MEAN_PRICE
## 1913  7      P4 2015    10   3        108    492.15      497    493.7667
## 1914  7      P7 2015     1   3         915    703.12      799    771.8340
## 1915  7      P3 2015     9  12          1    1002.32     1499   1188.8814
## 1916  7      P2 2015     4   4         131    695.13      879    790.7672
## 1917  7      P5 2015     7   4          23    797.05      799    798.1750
## 1918  7      P3 2015     8   8           3    1099.00     1499   1213.2857
##      SD_PRICE MEDIAN_PRICE WEEK MEAN_WD    DIF_WD    DIF_WEEK
## 1913    2.3879749      492.150    40   738.63  0.6663011    12.30375
## 1914   38.6591560      799.000     1   738.63  1.0817324   799.00000
## 1915  124.3456555     1185.575    37   738.63  1.6051000    32.04257
## 1916   59.9173717      760.710    14   738.63  1.0298932    54.33643
## 1917    0.7030776      798.000    27   738.63  1.0803785    29.55556
## 1918  184.0174825     1099.000    32   738.63  1.4878897    34.34375
```

Let's see how our model goes now.

```
fields <- setdiff(setdiff(names(train), "QTY_ORDER"), "MEDIAN_PRICE")
```

```
merged_data.hex <-
```

```
  as.h2o(merged_data, destination_frame = "merged_data2.hex")
```

```
##
|
|
|
|=====| 100%
```

```
data_to_model <-
```

```
  h2o.splitFrame(data = merged_data.hex ,
                 ratios = 0.80,
                 seed = 1234)
```

```
train <- data_to_model[[1]]
```

```
test <- data_to_model[[2]]
```

```
RF3 <- h2o.randomForest(
  x = fields,
  y = "QTY_ORDER",
  training_frame = train,
  validation_frame = test,
  ntrees = 150,
  max_depth = 45,
  seed = 1234
)
```

```
## Warning in .h2o.startModelJob(algo, params, h2oRestApiVersion): Dropping constant columns: [YEAR].
```

```
##
|
|
|
|===| 5%
|
```

```

|=====| 44%
|=====| 67%
|=====| 85%
|=====| 100%

```

```
h2o.r2(RF3, valid = T)
```

```
## [1] 0.4717004
```

We could get a precision of 47% this time. Let's see how this new tuning changed the model.

```
## Model Details:
```

```
## =====
```

```
##
```

```
## H2ORegressionModel: drf
```

```
## Model Key: DRF_model_R_1477243664145_254
```

```
## Model Summary:
```

```
##   number_of_trees number_of_internal_trees model_size_in_bytes min_depth
```

```
## 1           150                150          1381129           19
```

```
##   max_depth mean_depth min_leaves max_leaves mean_leaves
```

```
## 1          33  23.55333         273         870   727.82000
```

```
##
```

```
## H2ORegressionMetrics: drf
```

```
## ** Reported on training data. **
```

```
## ** Metrics reported on Out-Of-Bag training samples **
```

```
##
```

```
## MSE: 313319.3
```

```
## RMSE: 559.7493
```

```
## MAE: 213.1122
```

```
## RMSLE: 1.162761
```

```
## Mean Residual Deviance : 313319.3
```

```
##
```

```
##
```

```
## H2ORegressionMetrics: drf
```

```
## ** Reported on validation data. **
```

```
##
```

```
## MSE: 295811.9
```

```
## RMSE: 543.8859
```

```
## MAE: 228.3817
```

```
## RMSLE: 1.12914
```

```
## Mean Residual Deviance : 295811.9
```

```
##
```

```
##
```

```
##
```

```
##
```

```
## Scoring History:
```

```
##   timestamp      duration number_of_trees training_rmse
```

```
## 1 2016-10-23 21:00:03 0.007 sec           0
```

```
## 2 2016-10-23 21:00:03 0.023 sec           1    890.97959
```

```
## 3 2016-10-23 21:00:03 0.037 sec           2    788.72355
```

```
## 4 2016-10-23 21:00:03 0.053 sec           3    728.94894
```

```
## 5 2016-10-23 21:00:03 0.056 sec           4    679.64253
```

```
##   training_mae training_deviance validation_rmse validation_mae
```

```


## 1
## 2    291.50841    793844.62158    712.97027    258.89478
## 3    274.73954    622084.83473    581.06056    221.52541
## 4    254.60491    531366.55673    572.85010    221.30372
## 5    248.85136    461913.97013    577.91312    222.56125
## validation_deviance
## 1
## 2    508326.60115
## 3    337631.37608
## 4    328157.23671
## 5    333983.57166
##
## ---
##          timestamp    duration number_of_trees training_rmse
## 146 2016-10-23 21:00:07  3.830 sec           145    560.17583
## 147 2016-10-23 21:00:07  3.863 sec           146    559.69524
## 148 2016-10-23 21:00:07  3.898 sec           147    559.57692
## 149 2016-10-23 21:00:07  3.934 sec           148    559.63236
## 150 2016-10-23 21:00:07  3.969 sec           149    559.67846
## 151 2016-10-23 21:00:07  4.005 sec           150    559.74930
## training_mae training_deviance validation_rmse validation_mae
## 146    213.36352    313796.96165    543.69361    228.22117
## 147    213.30710    313258.76463    543.99337    228.32889
## 148    213.15701    313126.32922    543.58501    228.25758
## 149    213.09555    313188.37461    543.20822    227.84803
## 150    213.06989    313239.97713    543.55318    228.01741
## 151    213.11223    313319.28132    543.88589    228.38168
## validation_deviance
## 146    295602.73969
## 147    295928.78221
## 148    295484.66453
## 149    295075.16945
## 150    295450.05837
## 151    295811.86385
##
## Variable Importances: (Extract with `h2o.varimp`)
## =====
##
## Variable Importances:
##      variable relative_importance scaled_importance percentage
## 1  PROD_ID  16743229440.000000    1.000000    0.226808
## 2  MIN_PRICE  14175352832.000000    0.846632    0.192023
## 3  MEAN_PRICE  10871889920.000000    0.649331    0.147273
## 4    DAY  10690825216.000000    0.638516    0.144820
## 5  SD_PRICE   7796886016.000000    0.465674    0.105618
## 6  MAX_PRICE   7446653952.000000    0.444756    0.100874
## 7   MONTH   6096412160.000000    0.364112    0.082583

```

```

ggplot(data=F3_summary, aes(x = variable, y = scaled_importance, fill = variable)) +
  geom_bar(stat="identity", position = "dodge")

```



There were some increasing in importance for the fields, specially the MIN\_PRICE. Now our model had a good precision gain.

## Conclusions

Our model could predict the prices with a 47% score. To further improvements we could extract more relationships with the time period and price practiced, as it appears to have significant relationship